

ATHAR SAYED

PROFESSIONAL SUMMARY

Generative AI Engineer with 2+ years of **software engineering** experience, specializing in enterprise LLM applications, RAG pipelines, Agentic AI, Azure OpenAI, Python, and AI/LLMOps. Experienced in building production knowledge-indexing pipelines, AI applications, automated CI/CD workflows, and cloud-based GenAI infrastructure.

TECHNICAL SKILLS

- **Programming Languages:** Python | C++ | JavaScript
- **Generative AI & LLM:** LLMs | Retrieval-Augmented Generation (RAG) | LangChain | Agentic AI | Prompt Engineering | LLM Evaluation | OpenAI API | Embeddings | Natural Language Processing (NLP) | LLaMA | Model Context Protocol
- **ML & Data Frameworks:** TensorFlow | PyTorch | Keras | Apache Spark | Apache Kafka
- **Backend & APIs:** FastAPI | Flask | Streamlit | Gradio | Node.js | REST APIs | Client-Server Architecture
- **Cloud & AI Infrastructure:** Microsoft Azure | Azure OpenAI | Azure Blob Storage | Azure Automation Runbooks | Hybrid Workers | Google Cloud Platform (GCP) | Oracle Cloud Infrastructure (OCI)
- **DevOps & MLOps:** Docker | Kubernetes | Jenkins | Git | CI/CD | Prometheus | Grafana | GitHub Actions | Linux | Bash | OpenTelemetry
- **Databases:** MongoDB | MySQL | SQLite

EDUCATION

Mukesh Patel School of Technology, Management & Engineering, NMIMS	2024 - 2026
• M. Tech. – Artificial Intelligence (CGPA: 9.08/10)	
Symbiosis Institute of Engineering and Technology, Symbiosis International University	2019 - 2023
• B.Tech. – Electronics & Telecommunication Engineering	
• Honors in Mechatronics & Automation (CGPA: 7.11/10)	

PROFESSIONAL EXPERIENCE

Diebold Nixdorf, Mumbai, India

- Generative AI Intern** **February 2026 - Present**
- Engineered and optimized **Retrieval-Augmented Generation (RAG) pipelines** by refactoring Python preprocessing and indexing components, removing legacy dependencies, and restructuring workflows to improve retrieval quality, modularity, maintainability, and execution efficiency.
 - Architected and migrated **knowledge-indexing workflows to Azure OpenAI**, replacing local embedding infrastructure with the serverless **text-embedding-3-large** model and integrating cloud-based embedding generation into the end-to-end data pipeline.
 - Developed flexible **FAISS indexing architecture** supporting incremental updates, multiple agent modes, and dynamic raw-index selection, enabling multiple indexing strategies through a configurable and reusable pipeline.
 - Re-engineered the **Confluence data-processing architecture** to execute export, preprocessing, embedding generation, and index creation directly on **Jenkins agents**, eliminating dedicated compute infrastructure and simplifying the end-to-end knowledge ingestion workflow.
 - Built automated **Jenkins CI/CD workflows** for full and incremental knowledge-base generation, including Docker image build and validation, artifact packaging, timestamped versioning, and structured **GAV-based publishing to Nexus Repository**.
 - Engineered production **Jira Service Desk (JSD) indexing pipelines** supporting open, closed, and combined ticket modes, implementing separate data/KB-summary handling, empty-export resilience, and fixes across production and non-production workflows.
 - Rebuilt and stabilized **Azure Automation Runbook and Hybrid Worker infrastructure**, resolving VM execution failures through resource reallocation, dependency configuration, environment synchronization, and end-to-end validation of automated pipeline execution.
 - Re-engineered a Streamlit chatbot analytics application into a **React + Express** dashboard, replacing full in-memory loading of **~38K PostgreSQL** records with on-demand API-driven retrieval, implemented server-side pagination, filtering, sorting, and search, plus **LRU/TTL caching**, request coalescing, and PostgreSQL connection pooling.

FortytwoLabs, Pune, India

- Software Engineer** **June 2023 - July 2024**
- Managed build & release processes for 6+ production C++ libraries, ensuring stable **cross-platform deployment** (Linux, Windows, macOS).
 - Implemented cryptographic modules, including **policy signature** generation and verification, to enhance application security and improve cryptographic processing efficiency.
 - Developed and maintained **Makefile**-based build scripts to automate compilation, dependency management, and packaging for Linux environments.
 - Implemented application monitoring using **Prometheus** and developed **Grafana** dashboards to visualize application health, resource utilization, and operational metrics.
 - Conducted a Proof of Concept (PoC) using **OpenTelemetry** to evaluate distributed tracing capabilities and improve observability across application workflows.

FortytwoLabs, Pune, India

- Technical Intern** **Jan 2023 - May 2023**
- Built hands-on expertise in **Linux/CentOS administration**, including SSH/PuTTY, Bash scripting, filesystem/permissions management, troubleshooting, and server configuration.
 - Worked with **Git/GitLab**, **REST APIs**, **Postman**, **Docker**, and Linux-based application deployment; configured **Apache HTTPD** and **Nginx reverse proxy** setups.
 - Configured **MongoDB** on CentOS, performed CRUD operations and authentication, and implemented a **3-node replica set** with primary/secondary configuration and replication monitoring.
 - Developed and tested C++ **components** using Google Test and implemented **RSA**, **Shamir Secret Sharing**, **SHA-256/OpenSSL**, and C++-to-**WebAssembly (Emscripten)** workflows.

CERTIFICATIONS

- | | |
|--|-------------|
| • Oracle Cloud Infrastructure 2025 Certified AI Foundations Associate - Oracle | 2025 |
| • Oracle Cloud Infrastructure 2025 Certified Generative AI Professional - Oracle | 2025 |

PUBLICATION

- | | |
|--|--------------------|
| • IOT Based Crowd Detection and Stampede Avoidance using Predictive Analysis - IEEE Xplore | Apr 4, 2023 |
| • Predictive Analysis of Stress Based on Vital Physiological Metrics - Springer Nature | 2026 |